

TECHNICAL UNIVERSITY OF DENMARK

HIGH-PERFORMANCE COMPUTING

COURSE 02614

Assignment 2

Authors:

Mikkel JENSEN, s123184

Collaborator:

Oskar HINT, s161559

January 20, 2017



Contents

1	Introduction	1
2	Assignment	1
2.1	Matrix matrix multiplication problem	1
2.2	Poisson problem	10
2.3	Mutli-GPU	13

CUDA Driver Version / Runtime Version	8.0 / 8.0
CUDA Capability Major/Minor version number:	3.5
Total amount of global memory:	11439 MBytes (11995054080 bytes)
(15) Multiprocessors, (192) CUDA Cores/MP:	2880 CUDA Cores
GPU Max Clock rate:	745 MHz (0.75 GHz)
Memory Clock rate:	3004 Mhz
Memory Bus Width:	384-bit
L2 Cache Size:	1572864 bytes
Total amount of constant memory:	65536 bytes
Total amount of shared memory per block:	49152 bytes
Total number of registers available per block:	65536
Warp size:	32
Maximum number of threads per multiprocessor:	2048
Maximum number of threads per block:	1024
Max dimension size of a thread block (x,y,z):	(1024, 1024, 64)
Max dimension size of a grid size (x,y,z):	(2147483647, 65535, 65535)
Maximum memory pitch:	2147483647 bytes
Texture alignment:	512 bytes
Concurrent copy and kernel execution:	Yes with 2 copy engine(s)
Run time limit on kernels:	No
Integrated GPU sharing Host Memory:	No
Support host page-locked memory mapping:	Yes
Alignment requirement for Surfaces:	Yes
Device has ECC support:	Enabled
Device supports Unified Addressing (UVA):	Yes
Device PCI Domain ID / Bus ID / location ID:	0 / 132 / 0

Table 1: Selected specifications of the Tesla K40c GPU.

1 Introduction

In this report we investigate the performance of GPU's on matrix matrix multiplication from project 1 and the Poisson-problem from project 2. The two tasks are compute bound and memory bound respectively, so they are a good basis for exploring the capabilities of GPU acceleration on different types of problems. All GPU kernels are executed on the cluster Tesla K40c with selected specifications listed in table 1

2 Assignment

2.1 Matrix matrix multiplication problem

Matrix matrix multiplication was covered in project 1, and the operation requires $\sim N^3$ floating point operations performed on $\sim N^2$ elements, which means that the problem is compute bound for large matrices. This should be optimal for GPU execution due to the very high number of cores working in parallel. In the following, a GPU program is developed in CUDA 8.0, with incremental steps that would improve the performance.

2.1.1 GPU1

The first version is a sequential version only to be run on a single CUDA core. Performance wise, this is very uninteresting, but it is a good exercise to ensure that the framework of `cudaMalloc` and `cudaMemcpy` is set-up correctly. The code is shown below.

```
__host__
void matmult_gpu1(int m, int n, int k, double * A, double * B,
    double * C){

    double *d_A, *d_B, *d_C;
    cudaMalloc(&d_A, n*m*sizeof(double));
    cudaMalloc(&d_B, k*m*sizeof(double));
    cudaMalloc(&d_C, n*k*sizeof(double));

    cudaMemcpy(d_A, A, n*m*sizeof(double), cudaMemcpyHostToDevice);
    cudaMemcpy(d_B, B, k*m*sizeof(double), cudaMemcpyHostToDevice);

    cudaSeq<<<1,1>>>(m, n, k, d_A, d_B, d_C);
    cudaMemcpy(C, d_C, n*k*sizeof(double), cudaMemcpyDeviceToHost);

    cudaFree(d_A);
    cudaFree(d_B);
    cudaFree(d_C);
}
```

Listing 1: `gpu1` function with the `cuda` framework. The kernel `cudaSeq` is identical to the `mkn` permutation from project 1.

The kernel is called with `<<<1,1>>>` which signifies one block consisting of one thread, and thus the entire operation is done in sequential mode. The kernel is identical to the `mkn` permutation of project 1, and is not shown here. The framework shown above is identical in all subsequent functions. The performance of `gpu1` is compared to the DGEMM `cblas` library function run on 4 threads in figure 1. In this interval with small matrices the DGEMM routine performs at a quite constant level at around 5-10 GFlops/s, whereas `gpu1` does about 5 MFlops/s, and thus the performance is different by a factor of 1000. That is expected because a single GPU thread can never realize the parallelization potential of both the problem and the GPU.

2.1.2 GPU2

Obviously, using only 1 thread is not the way to use a GPU to increase performance, so in order to better utilize the many threads allowed on a GPU, we now implement a function, `gpu2`, that uses one thread per element of matrix C is size $m \times n$. The threads are ordered in a 2D grid consisting of 16×16 threads. The grid size is adapted to the matrix size such that the grid dimension are `gridx = ceil(n/16)` and `gridy = ceil(m/16)`. If the matrix size is not a multiple of the block size, excess threads are allocated and not used, but for large matrices, the number of unused threads is small compared to the total number of threads. For small matrices it is thus favourable to have a smaller grid size. The code is shown below.




Figure 1: The gpu1 function compared to the cblas DGEMM run on 4 cores.

```

__host__
void matmult_gpu2(int m, int n, int k, double * A, double * B,
    double * C){
    int K = 32;
    int gridx = ceil(n*1.0/K);
    int gridy = ceil(m*1.0/K);
    ...
    cudaPar<<<dim3(gridx, gridy), dim3(K,K)>>>(m,n,k,d_A,d_B,d_C);
    ...
}

__global__
void cudaPar(int m, int n, int k, double * A, double * B, double *
    C){
    int i = blockIdx.x*blockDim.x + threadIdx.x;
    int j = blockIdx.y*blockDim.y + threadIdx.y;

    if(i < n && j < m){
        C[i + j*n] = 0;

        for(int l = 0; l < k; l++){
            C[j*n + i] += A[k*j + l]*B[l*n + i];
        }
    }
}

```

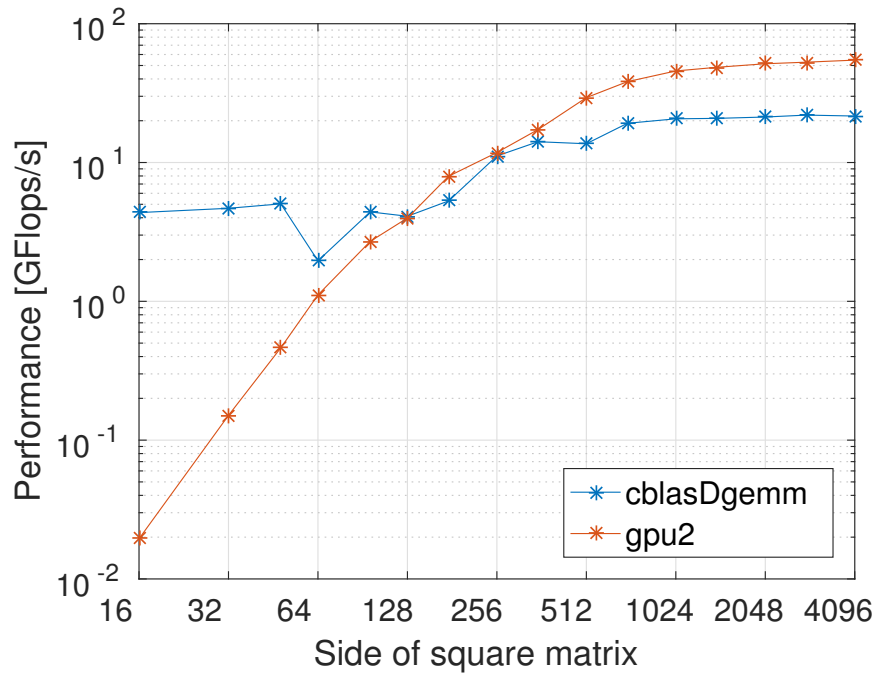


Figure 2: Performance of the `gpu2` function and the `cblas` DGEMM for different matrix sizes.

}

Listing 2: Code sample of the naive implementation with one thread per element in C .

The `if` clause ensures that work is only done inside the matrix, such that none of the excess threads do any unwanted work that could cause a segmentation fault. The last `for` loop runs through the k elements of A and B in row j and column i , and saves the answer sum of the products in the correct element of C .

The `gpu2` function is compared to the `cblas` DGEMM run on 4 cores for different matrix sizes in figure 2. For small matrix sizes, the `cblasDgemm` performs much better than `gpu2`. This is again due to the limited number of threads being able to operate on a small matrix, simply due to the limited number of elements. As the matrix size increases, the performance increases linearly in this loglog plot, which implies that the behaviour is potential, and most likely quadratic, due to the number of cores increasing with the number of elements, which is N^2 . At $N \approx 100$ the `gpu` version performs better, and for $N > 250$, the performance of `gpu2` flattens out, to stabilize at ~ 55 GFlops/s, which is around 2.5 times more than the `cblasDgemm`, which tops at ~ 22 GFlops/s. The flattening of the performance of `gpu2` is analysed with the Nvidia profiler. The results for square matrices of size 1024 show that L2 cache bandwidth is only used at $\sim 50\%$. Since the problem is compute bound, the way to improve performance is to increase the cache usage, and make use of the much higher bandwidth from cache to processor than from global memory to processor. When the cache memory is utilized satisfactorily, the memory usage from the global memory to the caches should be optimized. The matrix multiplication takes about 120 ms to run.

2.1.3 GPU3

As an attempt to improve memory usage from the caches the gpu function, we now changed `gpu2` to use half as many threads, but have each thread update 2 adjacent values of the C matrix. We choose to have each thread do its own element and the one below. This choice is made to ensure coalesced memory reads in matrix A , which is used row wise by each thread, and a single cache line of 128 bytes thus fits with the warp size of 32 threads. Matrix B is still read in a non-coalesced manner, but that will be the case no matter which neighbour is chosen. For larger matrices, decreasing the number of threads has no immediate effect on the performance due to the fact that the matrix multiplication of large matrices is a compute bound operation, the limiting factor is the available CUDA cores that can perform flops in parallel, but since we cannot change that we chose to improve the cache usage. The y -dimension of the grid is now `gridy = ceil(m/16/2)`, but otherwise the framework is the same. The kernel source code is seen below. The inner loop is over p to increase the memory re-usage of the B matrix, and it has been written out explicitly to get rid of the overhead. The first if clause takes care of the threads that can do both elements, and the second if clause take the threads which cannot compute both elements assigned to it due to the last one being out of bounds.

```
--global--
void cudaPar2(int m, int n, int k, int p, double * A, double * B,
    double * C){
    const int P = 2;
    double C_r[P]={0.0,0.0};

    int i = blockIdx.x*blockDim.x + threadIdx.x;
    int j = blockIdx.y*blockDim.y + threadIdx.y;

    j = j*p; // allows C[i + (0/1) j*n] indexing
    if(i < n && j < m-p + 1){
        for(int l = 0; l < k; l++){
            C_r[0] += A[(j + 0)*k + l]*B[n*l + i];
            C_r[1] += A[(j + 1)*k + l]*B[n*l + i];
        }
        for(int pp = 0; pp < p; pp++) C[(j + pp)*n + i] = C_r[pp];
    }
    if(i < n && j > m-p && j < m){
        for(int l = 0; l < k; l++){
            C_r[0] += A[(j + 0)*k + l]*B[n*l + i];
            C_r[1] += A[(j + 1)*k + l]*B[n*l + i];
        }
        for(int pp = 0; pp < p; pp++) C[(j + pp)*n + i] = C_r[pp];
    }
}
```

Listing 3: The source code for `textttgpu3`. The l loop is the outer loop to increase memory re-usage.

The performance of `gpu3` is compared to the `gpu2` and `cblasDgemm` in figure 3, and we see that `gpu3` performs better for $N > 100$, and increases to perform twice as well for a matrix




Figure 3: Performance of `gpu2`, `gpu3` and the `cblasDgemm` functions for increasing matrix sizes.

size of 4096 with ~ 100 GFlops/s. For a similar case analysed with `gpu2`, square matrices of size 1024 are investigated with `gpu3`. We have a 50 % increase in the usage of the L2 cache bandwidth and the global memory bandwidth, and as a result, the run time for each call is down by 17 % to about 100 ms.

2.1.4 GPU4

In an attempt to increase the reuse of the column of B each thread loads, we increased p to 4. For large matrices, there are still threads enough to have threads calculation on the CUDA cores, and the memory that is required to be loaded from the global memory is significantly reduced, which even though it is a compute bound problem, should increase the performance, due to better cache utilization. The source code for `gpu4` is much alike that of `gpu3`, but for completeness, it is included below.

```
--global--
void cudaPar4(int m, int n, int k, int p, double * A, double * B,
    double * C){
    const int P = 4;
    double C_r[P]={0.0,0.0,0.0,0.0};

    int i = blockIdx.x*blockDim.x + threadIdx.x;
    int j = blockIdx.y*blockDim.y + threadIdx.y;

    j = j*p; // allows C[i + (0/1) j*n] indexing
    if(i < n && j < m-p + 1){
```



```

    for(int l = 0; l < k; l++){
        C_r[0] += A[(j + 0)*k + l]*B[n*l + i];
        C_r[1] += A[(j + 1)*k + l]*B[n*l + i];
        C_r[2] += A[(j + 2)*k + l]*B[n*l + i];
        C_r[3] += A[(j + 3)*k + l]*B[n*l + i];
    }
    for(int pp = 0; pp < p; pp++) C[(j + pp)*n + i] = C_r[pp];
}
if(i < n && j > m-p && j < m){
    for(int l = 0; l < k; l++){
        C_r[0] += A[(j + 0)*k + l]*B[n*l + i];
        C_r[1] += A[(j + 1)*k + l]*B[n*l + i];
        C_r[2] += A[(j + 2)*k + l]*B[n*l + i];
        C_r[3] += A[(j + 3)*k + l]*B[n*l + i];
    }
    for(int pp = 0; pp < p; pp++) C[(j + pp)*n + i] = C_r[pp];
}
}

```

Listing 4: The source code for `textttgpu4`. The l loop is the outer loop to increase memory re-usage.

The performance of `gpu4` is compared to the previous functions in figure 4, and we observe a slight increase in performance to about 108 GFlops/s, but much the same behaviour as expected due to the similarity of the code. The profiler shows that the used bandwidth in both L2 cache and global memory is increased by $\sim 10\%$. The increase of p from two to four, only gave a 10 % increase, so increasing p further is not likely to be worthwhile.

2.1.5 GPU5

Instead, we now attempt to utilize the shared memory each block has. The bottleneck is now likely the non-coalesced reads from the B matrix, and so by utilising the shared memory one can read blocks of both matrices in a coalesced manner, store the loaded blocks in shared memory and do the multiplication on the blocks, store the result in a third shared memory variable, and then move the blocks along A and B , and at the end, the last shared memory variable can be saved to C . In this way, the L1 cache bandwidth is now utilised properly. The approach is similar to the blocking exercise in project 1, only here the memory management is done explicitly. The source code is shown below.

```

__global__
void cudaSMEM(int m, int n, int k, double * A, double * B, double
    * C){
    const int K = 32;
    int kk = k/K;

    int i = blockIdx.x*blockDim.x + threadIdx.x;
    int j = blockIdx.y*blockDim.y + threadIdx.y;
    int iBlock = threadIdx.x;
    int jBlock = threadIdx.y;

```




Figure 4: Performance of the gpu functions and the cblasDgemm as a function of matrix size.

```

__shared__ double smemA[K][K];
__shared__ double smemB[K][K];
__shared__ double smemC[K][K];

smemC[jBlock][iBlock] = 0;

for(int q = 0; q < kk; q++){
    smemB[jBlock][iBlock] = B[i + q*K*n + jBlock*n];
    smemA[jBlock][iBlock] = A[iBlock + j*k + q*K];
    __syncthreads();
    for(int z = 0; z < K; z++){
        smemC[jBlock][iBlock] += smemA[jBlock][z]*smemB[z][iBlock];
    }
    __syncthreads();
}
C[i + j*n] = smemC[jBlock][iBlock];
}

```

Listing 5: Source code for the shared memory kernel

In this case, a block size of 32 is used. This choice means that each warp now make up a single row in each block, and the cache lines consist of 32 floats, which means that each warp requires only 1 cache line per block, which optimizes the memory transfers. Each threads block does a matrix-block of the same size, so each thread does one element of C . Three 32

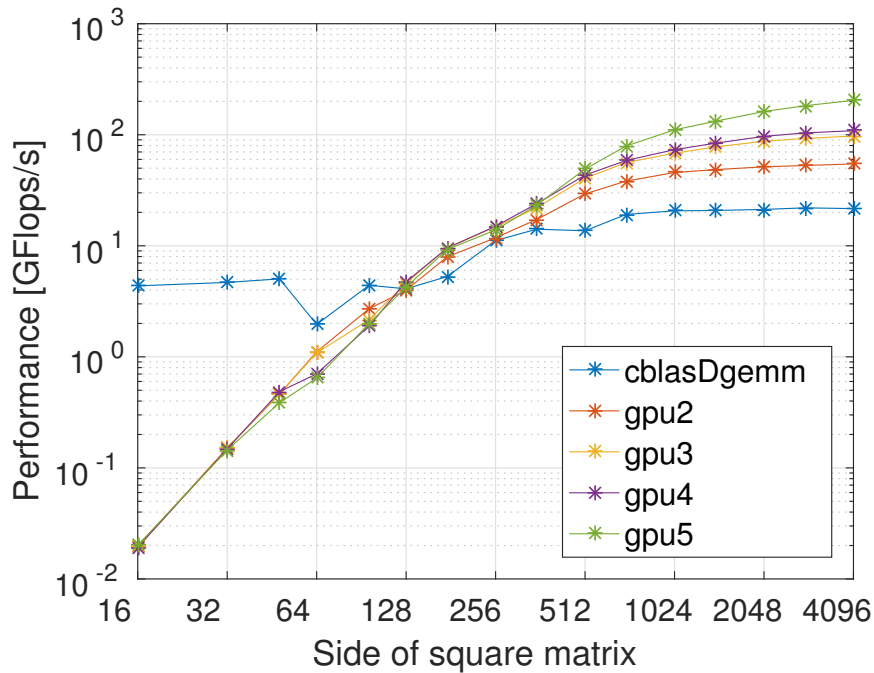


Figure 5: Performance of the gpu functions and the cblasDgemm library function. **gpu5** is seen to be roughly a factor of 2 faster than **gpu4**.

by 32 blocks is about 25 MB, so a block size of 64 will not fit in the shared memory which is 49 MB. Block sizes of non-multiples of 32 will cause unused data to be loaded, taking up bandwidth which is of course not desirable.

The performance of **gpu5** is added to the plot, and it is seen in figure 5. For large matrices of size $N > 2000$, the performance is about 200 GFlops/s, which is about twice as good as **gpu4**. Unfortunately, **gpu5** does not perform well on cuda 6.5, which is used for the analyser, so that can not back up the claims, but a factor of 2 speed up is a great result.

Compared to the theoretical maximum of $745 \text{ MHz} \times 2880 \approx 2 \text{ THz}$, we only utilize 10 % of the clock cycles, but as evident from figure 5, the **gpu5** curve is still increasing at $N = 4000$, so a higher fraction could be possible with larger matrices. The optimization is of course not finished with this, but no more optimizations are carried out. Further changes could be to split up C into smaller parts, to make use of the fact that the elements of C are independent, and thereby the problem allows for asynchronous memory copying. This could hide the latency, which for **gpu5** takes $\sim 10\%$ of the gpu time. Since the problem is compute bound, an obvious improvement would be to split up the matrix and compute each half on two different GPUs as it is done in the Poisson exercise. That would double the number of available CUDA cores, which should do great in bringing the wall time down, but in turn would not improve the resource utilization.

2.1.6 GPULIB

Finally, Nvidia has created a BLAS version called cublas optimized for the GPU, and as we implemented it as **gpulib**. The **cublasDgemm** takes the matrices as column major, so with the row major form of C , care must be taken to ensure a correct result. First, we realize that




Figure 6: Performance of gpu2 to gpu5, the cublas function, and the cblas function versus the matrix size.

a row major matrix read as column major is the transposed matrix, i.e.

$$A_C = A_R^T, \quad (1)$$

where superscript T denotes transpose, and subscript R/C denotes whether the matrix is read column wise or row wise. We assume that all matrices are stored row wise. We then realize that if $C_R = A_R \times B_R$, then $C_R^T = B_R^T \times A_R^T = B_C \times A_C = C_C$. So by B and A to cublasDgemm, the result $C_C = C_R^T$ is the transposed of what we are after. This, however, fixes itself, as C reads the matrix in row major format, effectively performing a transpose: $C_C^T = (C_R^T)^T = C_R$ which is the desired result.

The cublasDgemm routine is applied to the same matrix sizes used for the other functions, and the comparison can be seen in figure 6. The performance of the cublas function continues to increase at $N = 4096$, but there the performance is already 540 Gflops/s, which is about a quarter of the theoretical maximum. It is worth noting, that the cublas functions is the poorest of all for small matrices, highlighting the fact that it is optimized to large matrices. However, for small matrices with a side length of $N < 100$, one should never attempt to go to the GPU.

2.2 Poisson problem

In this section, the Poisson problem is revisited and implemented to run on a GPU. The Poisson problem is memory bound, and thus the limiting factor of the computation is the memory bandwidth to the cores. In all three subsequent subsections, the memory is initialized on the CPU and then transferred to the GPU for computations, and transferred back in the end.

2.2.1 Single thread on GPU

Again, we begin by implementing a kernel that runs exactly like the sequential CPU version, but on the GPU. This ensures that the cuda framework is correct, and that the kernel produces the right result. The framework is shown below, with the source code of `jacobi_seq_kernel` is identical to the sequential version in project 2.

```

cudaMalloc(&d_u, memsize);
cudaMalloc(&d_uo, memsize);
cudaMalloc(&d_f, memsize);
cudaMemcpy(d_u, u, memsize, cudaMemcpyHostToDevice);
cudaMemcpy(d_uo, uo, memsize, cudaMemcpyHostToDevice);
cudaMemcpy(d_f, f, memsize, cudaMemcpyHostToDevice);

int k = 0;
while(k < kmax){
    jacobi_seq_kernel<<<1, 1>>>(d_u, d_uo, d_f, N, delta2);
    double * temp = d_uo;
    d_uo = d_u;
    d_u = temp;
    k++;
}

cudaMemcpy(uo, d_uo, memsize, cudaMemcpyDeviceToHost);
cudaFree(d_u);
cudaFree(d_uo);
cudaFree(d_f);

```

In this project we do not consider the convergence, and thus, the program iterates until a user-specified maximum is reached. The timing of the function is measured for different matrix sizes ranging from 16×16 to 256×256 and compared to the `omp3` function from project 2. From project 2, we know that the optimal number of cores is dependent on the matrix size due to parallel overhead, so for matrices smaller than 100×100 , we will use a single core, and for matrices larger than that, we will use 2 cores. The resulting graph is shown in figure 7

2.2.2 Naive version with 1 thread per element

One thread to do all the work is obviously not good, so we implement a naive version of one thread per element of the matrices. Per iteration, each thread loads 5 values (4 neighbours from u_o and 1 from f) and performs 6 Flops, and the problem is thus memory bound. The timings are compared to the best OMP version in figure 8. For the OMP version, we use increasing number of threads with increasing matrix sizes.

SOMETHING ABOUT PROFILING OF SIN HERE. THE PROFILING NODE WAS/IS OVERUSED For improvements, an obvious choice would be to implement the shared memory used in `gpu5` is the `matmult` exercise, but since each point is only used by its neighbours, most points will be used inside a single block, enabling the automatic cache




Figure 7




Figure 8

control to do quite well. For the blocking to be really effective, the blocks should be iterated multiple times between being synchronized. If the blocks do e.g. 10 iterations isolated from the rest of the blocks, the flops/element would increase by a factor of 10 also. The number of iterations should be chosen to match the optimum match between flops and memory for the GPU used. This procedure would require additional calculations to get the edges of the blocks correct, but it would still be beneficial. Another optimization would be to have the loop and the pointer swap on the GPU. That way, each block could have a copy of its relevant region of f in the cache all the time, which would reduce the memory transfer from the global memory by a third, leaving more bandwidth for transferring u and u_o .

2.3 Mutli-GPU